

Topic Models

June 9, 2025

1 Assignment 5.1: Topic Modeling

Course: ADS 509 - Applied Text Mining

Assignment: Topic Modeling with NMF, LSA, and LDA

Student: Gabriel Elohi

Date: \$(date)

1.1 Overview

In this assignment, we will build and compare three different topic modeling approaches: 1. **NMF (Non-negative Matrix Factorization)** model 2. **LSA (Latent Semantic Analysis)** model 3. **LDA (Latent Dirichlet Allocation)** model

We will work with the Brown University corpus from NLTK and compare the resulting topic allocations with the official document classifications.

1.2 AI Tool Attribution

If any AI tools (ChatGPT, Gemini, GitHub Copilot, etc.) are used in this assignment, they will be explicitly disclosed and cited here with explanations of their contributions.

1.3 1. Setup and Data Exploration

Let's start by importing the necessary libraries and exploring the Brown corpus.

```
[1]: # Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter, defaultdict
import warnings
warnings.filterwarnings('ignore')

# NLTK imports
import nltk
from nltk.corpus import brown, stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
```

```

from nltk.tag import pos_tag

# Scikit-learn imports for topic modeling
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.decomposition import NMF, TruncatedSVD, LatentDirichletAllocation
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.preprocessing import normalize

# Download required NLTK data
nltk.download('brown', quiet=True)
nltk.download('stopwords', quiet=True)
nltk.download('punkt', quiet=True)
nltk.download('wordnet', quiet=True)
nltk.download('averaged_perceptron_tagger', quiet=True)

print("Libraries imported successfully!")

```

Libraries imported successfully!

1.3.1 1.1 Exploring the Brown Corpus

```

[2]: # Explore the Brown corpus structure
print("Brown Corpus Categories:")
categories = brown.categories()
print(f"Number of categories: {len(categories)}")
print(f"Categories: {categories}")

print("\nDocument counts per category:")
for category in categories:
    file_count = len(brown.fileids(categories=category))
    print(f"{category}: {file_count} documents")

```

Brown Corpus Categories:

Number of categories: 15

Categories: ['adventure', 'belles_lettres', 'editorial', 'fiction', 'government', 'hobbies', 'humor', 'learned', 'lore', 'mystery', 'news', 'religion', 'reviews', 'romance', 'science_fiction']

Document counts per category:

adventure: 29 documents

belles_lettres: 75 documents

editorial: 27 documents

fiction: 29 documents

government: 30 documents

hobbies: 36 documents

humor: 9 documents

learned: 80 documents

lore: 48 documents

```
mystery: 24 documents
news: 44 documents
religion: 17 documents
reviews: 17 documents
romance: 29 documents
science_fiction: 6 documents
```

```
[3]: # Get sample documents from different categories
print("Sample documents from different categories:")
for i, category in enumerate(categories[:3]):
    file_id = brown.fileids(categories=category)[0]
    sample_text = ' '.join(brown.words(file_id)[:50])
    print(f"\n{category.upper()} - {file_id}:")
    print(sample_text + "...")
```

Sample documents from different categories:

ADVENTURE - cn01:

Dan Morgan told himself he would forget Ann Turner . He was well rid of her . He certainly didn't want a wife who was fickle as Ann . If he had married her , he'd have been asking for trouble . But all of this was rationalization . Sometimes...

BELLES_LETTRES - cg01:

Northern liberals are the chief supporters of civil rights and of integration . They have also led the nation in the direction of a welfare state . And both in their objectives of non-discrimination and of social progress they have had ranged against them the Southerners who are called Bourbons...

EDITORIAL - cb01:

Assembly session brought much good The General Assembly , which adjourns today , has performed in an atmosphere of crisis and struggle from the day it convened . It was faced immediately with a showdown on the schools , an issue which was met squarely in conjunction with the governor...

1.3.2 1.2 Data Preprocessing

```
[4]: # Initialize preprocessing tools
lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    """
    Preprocess text by tokenizing, removing stopwords, and lemmatizing.
    """
    # Convert to lowercase and tokenize
    tokens = word_tokenize(text.lower())
```

```

    # Remove non-alphabetic tokens and stopwords
    tokens = [token for token in tokens if token.isalpha() and token not in stop_words]

    # Lemmatize tokens
    tokens = [lemmatizer.lemmatize(token) for token in tokens]

    # Filter out very short tokens
    tokens = [token for token in tokens if len(token) > 2]

    return ' '.join(tokens)

print("Preprocessing function defined.")

```

Preprocessing function defined.

```

[5]: # Prepare the corpus for topic modeling
print("Preparing corpus for topic modeling...")

documents = []
document_categories = []
document_ids = []

# Process documents from each category
for category in categories:
    file_ids = brown.fileids(categories=category)
    for file_id in file_ids:
        # Get raw text
        raw_text = ' '.join(brown.words(file_id))

        # Preprocess text
        processed_text = preprocess_text(raw_text)

        # Only include documents with sufficient content
        if len(processed_text.split()) > 50:
            documents.append(processed_text)
            document_categories.append(category)
            document_ids.append(file_id)

print(f"Total documents prepared: {len(documents)}")
print(f"Average document length: {np.mean([len(doc.split()) for doc in documents]):.1f} words")

```

Preparing corpus for topic modeling..

Total documents prepared: 500

Average document length: 1021.7 words

1.4 2. NMF (Non-negative Matrix Factorization) Topic Model

Let's start with building an NMF model for topic modeling.

```
[6]: # Create TF-IDF vectorizer for NMF
print("Creating TF-IDF matrix for NMF...")

# Parameters for vectorization
max_features = 1000 # Limit vocabulary size
min_df = 2 # Ignore terms that appear in less than 2 documents
max_df = 0.8 # Ignore terms that appear in more than 80% of documents

tfidf_vectorizer = TfidfVectorizer(
    max_features=max_features,
    min_df=min_df,
    max_df=max_df,
    ngram_range=(1, 2), # Include unigrams and bigrams
    stop_words='english'
)

tfidf_matrix = tfidf_vectorizer.fit_transform(documents)
feature_names = tfidf_vectorizer.get_feature_names_out()

print(f"TF-IDF matrix shape: {tfidf_matrix.shape}")
print(f"Vocabulary size: {len(feature_names)}")
```

```
Creating TF-IDF matrix for NMF...
TF-IDF matrix shape: (500, 1000)
Vocabulary size: 1000
```

```
[7]: # Check scikit-learn version for NMF parameter compatibility
import sklearn
print(f"Scikit-learn version: {sklearn.__version__}")

# Note: NMF parameters changed in different sklearn versions:
# - Older versions used 'alpha' parameter
# - Newer versions use 'alpha_W' and 'alpha_H' parameters
# We'll use basic parameters for maximum compatibility
```

```
Scikit-learn version: 1.6.1
```

```
[8]: # Fit NMF model
print("Fitting NMF model...")

n_topics = 10 # Number of topics to extract
random_state = 42

# Create NMF model with compatible parameters
nmf_model = NMF(
```

```

    n_components=n_topics,
    random_state=random_state,
    init='nndsvd', # Non-negative double SVD initialization
    solver='cd', # Coordinate descent solver
    max_iter=200
)

nmf_topics = nmf_model.fit_transform(tfidf_matrix)

print(f"NMF model fitted successfully!")
print(f"Document-topic matrix shape: {nmf_topics.shape}")

```

Fitting NMF model...

NMF model fitted successfully!

Document-topic matrix shape: (500, 10)

```

[9]: # Display top words for each NMF topic
def display_topics(model, feature_names, n_top_words=10, model_name="Model"):
    """
    Display the top words for each topic.
    """
    print(f"\n=== {model_name} Topics ===")

    for topic_idx, topic in enumerate(model.components_):
        top_words_idx = topic.argsort()[::-1][:n_top_words]
        top_words = [feature_names[i] for i in top_words_idx]
        top_weights = [topic[i] for i in top_words_idx]

        print(f"\nTopic {topic_idx + 1}:")
        for word, weight in zip(top_words, top_weights):
            print(f"    {word}: {weight:.3f}")

# Display NMF topics
display_topics(nmf_model, feature_names, n_top_words=8, model_name="NMF")

```

=== NMF Topics ===

Topic 1:

```

said: 1.231
like: 0.477
got: 0.395
know: 0.384
woman: 0.352
went: 0.343
room: 0.341
thought: 0.337

```

Topic 2:

state: 0.450
nation: 0.422
united: 0.370
war: 0.366
soviet: 0.357
american: 0.332
government: 0.323
united state: 0.298

Topic 3:

human: 0.264
man: 0.245
life: 0.241
experience: 0.231
literature: 0.217
idea: 0.204
form: 0.203
century: 0.203

Topic 4:

temperature: 0.381
surface: 0.353
used: 0.285
cell: 0.274
data: 0.237
pressure: 0.231
material: 0.230
fig: 0.222

Topic 5:

school: 0.845
student: 0.610
college: 0.593
child: 0.407
education: 0.320
teacher: 0.305
university: 0.278
class: 0.186

Topic 6:

music: 0.495
song: 0.340
new: 0.303
performance: 0.266
new york: 0.258
york: 0.254
miss: 0.254

jazz: 0.254

Topic 7:

church: 1.098
god: 0.490
christian: 0.388
catholic: 0.298
christ: 0.293
john: 0.218
member: 0.184
new: 0.158

Topic 8:

man: 0.383
men: 0.354
horse: 0.321
water: 0.309
head: 0.298
foot: 0.293
eye: 0.292
gun: 0.290

Topic 9:

cost: 0.348
company: 0.325
sale: 0.310
industry: 0.279
program: 0.250
plant: 0.249
market: 0.246
tax: 0.241

Topic 10:

president: 0.698
kennedy: 0.458
said: 0.416
state: 0.360
committee: 0.344
house: 0.326
election: 0.280
party: 0.261

1.4.1 2.1 NMF Topic Interpretation

Analysis of NMF Topics:

[Add your interpretation of the NMF topics here. Discuss what themes or subjects each topic seems to represent based on the top words.]

1.5 3. LSA (Latent Semantic Analysis) Topic Model

Now let's build an LSA model using Truncated SVD.

```
[10]: # Fit LSA model using TruncatedSVD
print("Fitting LSA model...")

lsa_model = TruncatedSVD(
    n_components=n_topics,
    random_state=random_state,
    algorithm='randomized'
)

lsa_topics = lsa_model.fit_transform(tfidf_matrix)

print(f"LSA model fitted successfully!")
print(f"Document-topic matrix shape: {lsa_topics.shape}")
print(f"Explained variance ratio: {lsa_model.explained_variance_ratio_.sum():.
↪3f}")
```

Fitting LSA model...

LSA model fitted successfully!

Document-topic matrix shape: (500, 10)

Explained variance ratio: 0.155

```
[11]: # Display LSA topics
display_topics(lsa_model, feature_names, n_top_words=8, model_name="LSA")
```

=== LSA Topics ===

Topic 1:

said: 0.219
man: 0.135
like: 0.128
new: 0.112
year: 0.110
day: 0.097
state: 0.092
house: 0.089

Topic 2:

said: 0.235
like: 0.122
eye: 0.109
got: 0.106
door: 0.103
looked: 0.095
knew: 0.090

went: 0.088

Topic 3:

said: 0.240
state: 0.217
president: 0.165
government: 0.108
kennedy: 0.106
united: 0.101
house: 0.099
tax: 0.096

Topic 4:

church: 0.213
god: 0.135
student: 0.106
school: 0.106
college: 0.103
christian: 0.101
world: 0.093
life: 0.091

Topic 5:

school: 0.328
student: 0.222
college: 0.206
child: 0.164
miss: 0.155
university: 0.121
education: 0.108
year: 0.107

Topic 6:

church: 0.248
said: 0.237
school: 0.124
christian: 0.112
social: 0.110
god: 0.107
catholic: 0.102
law: 0.099

Topic 7:

church: 0.531
god: 0.232
john: 0.170
christian: 0.159
christ: 0.146

```
river: 0.124
catholic: 0.122
water: 0.109
```

Topic 8:

```
school: 0.266
college: 0.250
student: 0.249
child: 0.159
teacher: 0.132
education: 0.115
war: 0.105
president: 0.105
```

Topic 9:

```
church: 0.234
president: 0.185
kennedy: 0.165
soviet: 0.120
game: 0.119
miss: 0.102
room: 0.098
music: 0.096
```

Topic 10:

```
president: 0.198
cell: 0.178
john: 0.133
kennedy: 0.133
temperature: 0.129
trial: 0.114
fig: 0.111
house: 0.109
```

1.5.1 3.1 LSA Topic Interpretation

Analysis of LSA Topics:

[Add your interpretation of the LSA topics here. Compare with NMF results and discuss similarities/differences.]

1.6 4. LDA (Latent Dirichlet Allocation) Topic Model

Finally, let's build an LDA model. LDA works better with count data rather than TF-IDF.

```
[12]: # Create count vectorizer for LDA
print("Creating count matrix for LDA...")

count_vectorizer = CountVectorizer(
```

```

    max_features=max_features,
    min_df=min_df,
    max_df=max_df,
    ngram_range=(1, 1), # Only unigrams for LDA
    stop_words='english'
)

count_matrix = count_vectorizer.fit_transform(documents)
count_feature_names = count_vectorizer.get_feature_names_out()

print(f"Count matrix shape: {count_matrix.shape}")
print(f"Vocabulary size: {len(count_feature_names)}")

```

Creating count matrix for LDA...
Count matrix shape: (500, 1000)
Vocabulary size: 1000

```

[13]: # Note: LDA parameter names in scikit-learn:
# - doc_topic_prior (equivalent to alpha in other LDA implementations)
# - topic_word_prior (equivalent to beta in other LDA implementations)
print("Setting up LDA model with scikit-learn parameter names...")

```

Setting up LDA model with scikit-learn parameter names...

```

[14]: # Fit LDA model
print("Fitting LDA model...")

# Create LDA model with correct parameter names
lda_model = LatentDirichletAllocation(
    n_components=n_topics,
    random_state=random_state,
    doc_topic_prior=0.1, # Document-topic concentration (alpha)
    topic_word_prior=0.01, # Topic-word concentration (beta)
    max_iter=100,
    learning_method='batch'
)

lda_topics = lda_model.fit_transform(count_matrix)

print(f"LDA model fitted successfully!")
print(f"Document-topic matrix shape: {lda_topics.shape}")
print(f"Log likelihood: {lda_model.score(count_matrix):.2f}")

```

Fitting LDA model...
LDA model fitted successfully!
Document-topic matrix shape: (500, 10)
Log likelihood: -1349398.43

```
[15]: # Display LDA topics
display_topics(lda_model, count_feature_names, n_top_words=8, model_name="LDA")
```

```
=== LDA Topics ===
```

```
Topic 1:
```

```
new: 397.064
city: 310.589
south: 219.396
year: 204.915
town: 179.862
day: 173.985
north: 161.176
old: 141.456
```

```
Topic 2:
```

```
social: 310.731
law: 284.947
state: 282.828
people: 269.414
community: 215.618
fact: 210.616
policy: 210.547
question: 194.406
```

```
Topic 3:
```

```
number: 317.757
point: 218.305
line: 183.520
form: 183.381
value: 178.789
data: 157.376
information: 156.605
used: 149.535
```

```
Topic 4:
```

```
said: 1539.817
like: 973.071
man: 772.658
know: 565.163
hand: 536.752
little: 526.372
day: 507.702
came: 480.362
```

```
Topic 5:
```

```
life: 337.321
```

work: 334.181
new: 301.967
world: 258.195
great: 233.810
say: 229.262
man: 228.720
book: 224.913

Topic 6:

school: 630.635
child: 398.718
student: 311.352
year: 309.762
college: 298.311
university: 218.312
program: 184.637
education: 183.863

Topic 7:

state: 644.998
year: 474.620
cost: 357.573
program: 278.977
tax: 246.010
business: 244.003
service: 228.834
development: 227.840

Topic 8:

church: 428.934
god: 299.323
life: 158.465
john: 157.679
new: 148.876
christian: 146.015
man: 139.350
religious: 119.690

Topic 9:

state: 415.560
president: 360.202
new: 322.374
war: 321.515
american: 310.930
united: 268.857
said: 246.264
nation: 225.687

Topic 10:
 used: 249.765
 surface: 210.217
 water: 182.212
 material: 167.622
 area: 162.748
 use: 161.222
 temperature: 161.010
 high: 126.234

1.6.1 4.1 LDA Topic Interpretation

Analysis of LDA Topics:

[Add your interpretation of the LDA topics here. Compare with NMF and LSA results.]

1.7 5. Model Comparison and Analysis

Let's compare the three topic modeling approaches and analyze their performance.

```
[16]: # Create a comparison of topic assignments
def get_dominant_topic(doc_topic_matrix):
    """
    Get the dominant topic for each document.
    """
    return np.argmax(doc_topic_matrix, axis=1)

# Get dominant topics for each model
nmf_dominant_topics = get_dominant_topic(nmf_topics)
lsa_dominant_topics = get_dominant_topic(lsa_topics)
lda_dominant_topics = get_dominant_topic(lda_topics)

# Create comparison DataFrame
comparison_df = pd.DataFrame({
    'document_id': document_ids,
    'true_category': document_categories,
    'nmf_topic': nmf_dominant_topics,
    'lsa_topic': lsa_dominant_topics,
    'lda_topic': lda_dominant_topics
})

print("Topic assignment comparison:")
print(comparison_df.head(10))
```

Topic assignment comparison:

	document_id	true_category	nmf_topic	lsa_topic	lda_topic
0	cn01	adventure	0	0	3
1	cn02	adventure	7	0	3
2	cn03	adventure	7	0	3

3	cn04	adventure	7	0	3
4	cn05	adventure	0	0	3
5	cn06	adventure	0	0	3
6	cn07	adventure	7	0	3
7	cn08	adventure	7	0	3
8	cn09	adventure	0	0	3
9	cn10	adventure	0	0	3

```
[17]: # Analyze topic distribution by true categories
print("\nTopic distribution analysis:")

for model_name, topic_col in [('NMF', 'nmf_topic'), ('LSA', 'lsa_topic'),
                               ('LDA', 'lda_topic')]:
    print(f"\n{model_name} Topic Distribution by Category:")
    topic_category_crosstab = pd.crosstab(comparison_df['true_category'],
    ↪ comparison_df[topic_col])
    print(topic_category_crosstab)
```

Topic distribution analysis:

NMF Topic Distribution by Category:

nmf_topic	0	1	2	3	4	5	6	7	8	9
true_category										
adventure	12	0	0	0	0	0	0	17	0	0
belles_lettres	5	10	31	1	5	10	2	6	0	5
editorial	0	16	0	0	2	3	3	1	1	1
fiction	14	0	0	0	0	2	3	10	0	0
government	0	7	1	3	1	0	0	0	16	2
hobbies	0	1	0	12	4	5	0	4	10	0
humor	5	0	0	0	0	3	0	0	0	1
learned	0	4	17	38	6	2	0	1	9	3
lore	1	5	9	4	7	4	4	7	4	3
mystery	18	0	0	0	0	0	0	6	0	0
news	0	4	0	0	7	11	2	1	7	12
religion	0	0	5	0	0	0	11	1	0	0
reviews	0	2	1	1	0	13	0	0	0	0
romance	22	0	0	0	1	0	1	4	1	0
science_fiction	3	0	2	0	0	0	1	0	0	0

LSA Topic Distribution by Category:

lsa_topic	0	2	4	5	6	7	9
true_category							
adventure	29	0	0	0	0	0	0
belles_lettres	74	0	0	0	0	1	0
editorial	27	0	0	0	0	0	0
fiction	27	0	0	0	2	0	0
government	25	3	1	0	0	0	1

hobbies	34	0	1	0	0	0	1
humor	9	0	0	0	0	0	0
learned	62	1	4	2	0	1	10
lore	41	0	2	1	3	1	0
mystery	24	0	0	0	0	0	0
news	43	1	0	0	0	0	0
religion	14	0	0	0	3	0	0
reviews	17	0	0	0	0	0	0
romance	29	0	0	0	0	0	0
science_fiction	6	0	0	0	0	0	0

LDA Topic Distribution by Category:

lda_topic	0	1	2	3	4	5	6	7	8	9
true_category										
adventure	0	0	0	29	0	0	0	0	0	0
belles_lettres	3	16	2	15	30	0	0	0	9	0
editorial	0	0	0	3	4	0	1	0	19	0
fiction	0	0	0	28	1	0	0	0	0	0
government	0	3	1	0	0	1	18	0	5	2
hobbies	6	0	1	3	5	2	4	1	1	13
humor	0	0	0	8	1	0	0	0	0	0
learned	2	8	19	2	13	3	7	0	3	23
lore	4	8	0	11	8	6	2	1	5	3
mystery	0	0	0	24	0	0	0	0	0	0
news	10	1	0	4	1	8	7	0	13	0
religion	0	5	1	2	3	0	0	6	0	0
reviews	1	0	0	0	15	0	0	0	1	0
romance	0	0	0	29	0	0	0	0	0	0
science_fiction	0	0	0	6	0	0	0	0	0	0


```
[18]: # Visualize topic distributions
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

models = [('NMF', nmf_topics), ('LSA', lsa_topics), ('LDA', lda_topics)]

for idx, (model_name, topic_matrix) in enumerate(models):
    # Calculate average topic weights
    avg_topic_weights = np.mean(topic_matrix, axis=0)

    axes[idx].bar(range(len(avg_topic_weights)), avg_topic_weights)
    axes[idx].set_title(f'{model_name} Average Topic Weights')
    axes[idx].set_xlabel('Topic')
    axes[idx].set_ylabel('Average Weight')
    axes[idx].set_xticks(range(len(avg_topic_weights)))

plt.tight_layout()
plt.show()
```



1.8 6. Conclusions and Insights

1.8.1 6.1 Model Comparison Summary

NMF (Non-negative Matrix Factorization): - *[Add your analysis of NMF performance and characteristics]*

LSA (Latent Semantic Analysis): - *[Add your analysis of LSA performance and characteristics]*

LDA (Latent Dirichlet Allocation): - *[Add your analysis of LDA performance and characteristics]*

1.8.2 6.2 Comparison with Official Brown Corpus Categories

[Discuss how well each model's topics align with the official Brown corpus categories. Which model performed best at capturing the underlying document structure?]

1.8.3 6.3 Key Findings

1. *[Finding 1]*
2. *[Finding 2]*
3. *[Finding 3]*

1.8.4 6.4 Recommendations

[Based on your analysis, which topic modeling approach would you recommend for different use cases?]

1.9 7. Additional Analysis (Optional)

1.9.1 7.1 Topic Coherence Analysis

[If time permits, add topic coherence analysis or other advanced metrics]

```
[19]: # Optional: Add any additional analysis code here
print("Assignment completed successfully!")
print("\nNext steps:")
```

```
print("1. Fill in the interpretation sections with your analysis")
print("2. Run all cells and verify results")
print("3. Convert notebook to PDF for submission")
print("4. Commit and push to GitHub")
```

Assignment completed successfully!

Next steps:

1. Fill in the interpretation sections with your analysis
2. Run all cells and verify results
3. Convert notebook to PDF for submission
4. Commit and push to GitHub